

Contents

The Ideal Development Team at Semios

Vision

L2 Checklist

The Ideal Team by Dimension

Developer Role Profile

Tech Lead Profile

Role Appendix

1

1

1

2

4

5

5

The Ideal Development Team at Semios

Target Maturity: L2 — Integrated Baseline: March 2026 (Org Average: L1 / 1.25)

Vision

The Ideal Dev Team at Semios positions AI as a first-class shared team capability. Every developer works from the same context, uses the same tools with shared configuration, and produces code that passes automated AI quality gates before human review. Knowledge is documented in the codebase, not locked in individual heads. This is L2: **Integrated** in the Maturity Model.

L2 Checklist

Teams are measured against the **AI Maturity Scorecard** (4 dimensions, 12 sub-dimensions, L1-L4 scale). The Ideal Team achieves **L2 (1.5+) on all 12 sub-dimensions**.

Sub-dimension	L2 Requirement	Verified By
AI Tool Adoption	Team standardized on Claude Code + Copilot PR Review	Config audit
Prompt & Context Engineering	CLAUDE.md in every active repo	Repo check
Agent Configuration	At least 1 custom skill or Copilot PR instruction set	Config audit
CI/CD Integration	Automated AI review on every PR	CI pipeline check
Ticketing & Planning	AI used to validate or refine tickets	Team self-report
Cross-System Connectivity	2+ integrations active (JIRA + GitHub) via MCP or CLI	Config audit
Quality Controls	PR review checklist with AI-specific items	Process doc
Security & Compliance	Documented AI usage policy at team level	Policy doc

Sub-dimension	L2 Requirement	Verified By
Measurement & KPIs	Adoption rate + PR review time tracked	Dashboard
Ways of Working	“AI Ways of Working” doc exists and is current	Doc review
Accountability & Ownership	Named AI champion with explicit responsibilities	Team roster
Scalability & Knowledge Transfer	Onboarding materials include AI workflow guide	Onboarding check

Scoring

- **L2 achieved:** All 12 sub-dimensions at 1.5+
- **Strong L2:** 8+ sub-dimensions at 2.0+
- **L2 with gaps:** 9-11 sub-dimensions at 1.5+ (targeted remediation)
- **Not yet L2:** Fewer than 9 sub-dimensions at 1.5+

Progress is assessed at the **end of each sprint/cycle**.

The Ideal Team by Dimension

1. Capability

AI Tool Adoption

L2: Standardized on 1-2 tools; licenses managed

The team has converged on **Claude Code as the primary AI tool** with Copilot for PR reviews. Every developer uses the same core toolchain. There is no tool fragmentation — the team doesn’t split across Claude, Cursor, Gemini, and ChatGPT without coordination.

Prompt & Context Engineering

L2: Shared templates; implementation standards at repo level

Every repository has a **CLAUDE.md file** checked into the root. It contains: project architecture, key commands (build, test, lint), coding conventions, known gotchas, and domain context. Developers don’t rebuild context from scratch each session — they start with a shared baseline and extend it per task.

The team maintains **shared skills and commands** for recurring tasks: bug investigation, PR creation, test writing, code review. These are documented, not tribal.

Agent Configuration

L2: Slash commands and custom instructions configured

The team uses **slash commands** for common workflows (PR creation, bug triage, review). At least one **custom skill** exists for the team's most repeated task. Copilot instruction sets are configured for automated PR review feedback.

2. Integration

CI/CD Integration

L2: AI-generated code goes through automated AI checks

Every PR passes through at least **two AI review layers** before human review: a pre-push check (Claude /review) and an automated check at the PR level (Copilot auto-review). AI doesn't just generate code — it validates it before a human ever sees it.

Ticketing & Planning

L2: AI assists in writing or refining tickets

Developers use AI to **validate ticket quality before starting work** — checking for missing acceptance criteria, ambiguous requirements, or missing edge cases. PMs use AI to refine PRDs. Tickets reference relevant code paths and architecture context.

Cross-System Connectivity

L2: AI reads from repos and docs; limited write access

The team has **at least 2 MCP servers active** connecting AI to the systems they use daily (JIRA, GitHub, Confluence, Slack). Developers can pull ticket context, check PR status, or reference documentation without leaving their AI session.

3. Governance

Quality Controls

L2: Linting and checks on AI-generated code; review checklists exist

AI-generated code is held to the **same quality bar as human code** — plus additional scrutiny. The team has a documented **PR review checklist** that includes AI-specific items: utility duplication check, test coverage verification, and confirmation that the author understands the code they're submitting.

Security & Compliance

L2: AI usage policy exists; approved tools; data handling rules

The team operates under a **documented AI usage policy** that covers: approved tools, data handling (what can/cannot be sent to AI), review requirements for AI-generated code, and expectations for unsupervised AI usage. Developers know the boundaries.

Measurement & KPIs

L2: Basic metrics tracked; adoption rate and usage frequency reported

The team **tracks at minimum**: AI adoption rate (who is using which tools), PR review time (before/after AI review layers), and cycle time per ticket. These are reported, not for surveillance, but to demonstrate impact and identify where AI helps vs. hurts.

4. Execution Ownership

Ways of Working

L2: Shared conventions documented in wiki or README

The team has a **documented “AI Ways of Working” page** covering: how to start an AI session (open CLAUDE.md, load context), when to use plan mode vs. direct generation, PR submission expectations for AI-generated code, and the team’s review workflow. New team members can read this and be productive with AI on day one.

Accountability & Ownership

L2: Champion owns AI adoption for the team

One person on the team is the **designated AI champion** — responsible for maintaining CLAUDE.md, curating the prompt library, onboarding new team members to AI workflows, and serving as the team’s contact for cross-team pattern sharing. This is an explicit responsibility, not a side interest.

Scalability & Knowledge Transfer

L2: Documentation and onboarding materials exist

When a new developer joins the team, they receive: a CLAUDE.md that gives AI immediate code-base context, a prompt library for common tasks, a documented review workflow, and a 30-minute walkthrough from the AI champion. Knowledge lives in artifacts, not in heads.

Developer Role Profile

1. **Starts every session from CLAUDE.md** — never rebuilds context from scratch
2. **Uses plan mode for non-trivial work** — reviews the AI’s plan before accepting generated code
3. **Runs /review before pushing** — catches issues before they reach human reviewers
4. **Understands the code they submit** — can explain any AI-generated PR line-by-line; flags sections where confidence is low
5. **Contributes to shared context** — updates CLAUDE.md when they discover new conventions, gotchas, or domain knowledge
6. **Uses MCPs for cross-system context** — pulls ticket details, checks CI status, and references docs within AI sessions

7. **Writes tests alongside AI-generated code** — uses AI to generate tests but validates coverage against edge cases
-

Tech Lead Profile

1. **Maintains the team's CLAUDE.md** — ensures it's current, reviewed, and checked into every active repo
 2. **Owns the AI review workflow** — defines and enforces the team's multi-layer review process
 3. **Curates the prompt library** — collects proven prompts from team members and documents them for reuse
 4. **Tracks adoption** — knows who is using which tools, identifies training gaps, and reports on AI impact
 5. **Champions cross-team sharing** — brings innovations from other teams and shares the team's own breakthroughs
 6. **Sets AI usage norms** — establishes when AI is appropriate (green light), when extra review is needed (yellow light), and when AI should not be used (red light — e.g., sensitive data, compliance-critical paths)
-

Role Appendix

The following role-specific guidance extends the Ideal Team vision to non-engineering roles that participate in the development workflow.

Product Manager

- Uses AI to **validate and refine PRDs** before handoff to engineering
- Ensures tickets include sufficient context for both developers and AI (acceptance criteria, relevant code paths, domain context)
- Ties feature investment to **measurable outcomes** (ARR, adoption, engagement) — not just delivery
- Participates in a **PM guild** for horizontal alignment across product teams

QA

- Uses AI for **test case generation** targeting edge cases and negative scenarios, not just happy paths
- Operates under a **clear data privacy policy** specifying what client data can and cannot be sent to AI tools
- Uses AI for **PR explanation** — understanding AI-generated code changes before testing
- Contributes QA-specific context to CLAUDE.md (test commands, known flaky areas, data constraints)

UX Designer

- Maintains the **design system as a living artifact** — Figma components, Storybook implementation, and design tokens stay in sync
- Uses AI for **design-to-code translation** (Figma Make, v0) within an established SOP
- Has a **design review gate** in the workflow — product cannot bypass UX and send designs directly to dev
- Contributes design context to CLAUDE.md (component naming conventions, design tokens, accessibility requirements)